

Replayed real-LLM deceivers (4 prompts × 2 attacker models × 6 benchmarks)

