


E11. Classical crowdsourcing estimators collapse when liars are the majority; the anchor does not



(b) Per benchmark at f = 0.7 (7 of 10 agents lie)

Known-channel oracle	91	86	95	93	97	82	91
RACE (ours)	90	85	95	92	97	78	90
AIP gated	51	48	59	37	97	73	61
IWMV	22	23	18	31	5	7	18
MACE	22	23	18	31	5	7	18
GLAD	22	23	18	25	5	7	17
KOS (binary only)	n/a	n/a	n/a	25	n/a	n/a	n/a
Dawid-Skene (label-free)	22	23	18	31	5	7	18
Dawid-Skene full	22	23	18	31	n/a	n/a	n/a
Majority vote	28	28	22	38	11	10	23
Receiver alone	81	79	91	89	95	70	84
	MMLU	MedQA	ARC	BoolQ	GSM8K	MATH-500	Mean

Means over six benchmarks, six attack settings (coherent, gate-aware p = 0.25, four LLM prompts) and three seeds, TEST split. KOS is defined for binary questions only; full-confusion Dawid-Skene is not run on open-answer benchmarks (n/a).